

AI in Fraud Detection: A Case Study of PayPal

ITAI 2372 | Francisco Medina | June 6, 2026 | Anubis | Prof. Sridhar Ganti

PayPal Holdings, Inc. is one of the world's largest digital payment platforms, founded in 1998 and headquartered in San Jose, California. As of 2022, the company serves over 435 million active accounts across more than 200 markets worldwide. As shown in Figure 1, PayPal's total payment volume (TPV) grew from \$578 billion in 2018 to \$1.36 trillion in 2022, representing a compound annual growth rate of approximately 23.8% showing the enormous scale of transactions the platform needs to protect from fraud (PayPal Holdings, 2023).


Year	Total Payment Volume Growth	TPV	YoY Growth
2018		\$578B	—
2019		\$712B	+23.2%
2020		\$936B	+31.5%
2021		\$1.25T	+33.9%
2022		\$1.36T	+8.5%

Figure 1. PayPal Total Payment Volume Growth (2018–2022). The rapid scaling of transaction volume made AI-driven fraud detection operationally essential. Source: PayPal Annual Reports.

This growth trajectory is significant because it illustrates why static, rule-based fraud detection became unsustainable. When PayPal was processing \$578 billion a year, rule sets could still be managed. At \$1.36 trillion across 200+ markets, the sheer volume, speed, and geographic spread made that approach unworkable. Early in its history, PayPal's rule-based systems flagged transactions against static thresholds such as transaction size, geographic location, and purchase velocity. These systems produced enormous numbers of false positives and were reactive by design, unable to catch new fraud patterns until investigators had already documented them (Columbus, 2017).

To address these structural weaknesses, PayPal made a strategic investment in artificial intelligence, building one of the most sophisticated real-time fraud detection platforms in global fintech. The result is a fraud loss rate of approximately 0.32% of total payment volume, well below the industry average that frequently exceeds 1.00% (DigitalDefynd, 2024). The broader academic literature on AI in finance confirms this outcome is consistent with expectations: research shows that classification and early warning systems built on machine learning outperform traditional methods in both accuracy and efficiency, making fraud detection one of the highest-impact uses of AI in the financial sector (Bahoo et al., 2024). This report examines how PayPal built that system, covering the technologies used, the results it produced, and the challenges the company faced along the way.

PayPal's fraud detection infrastructure is a layered ensemble of AI systems, each contributing a specialized capability. Figure 2 illustrates the end-to-end detection pipeline through which every transaction is routed, evaluating over 700 distinct risk signals and returning a scored decision in under 300 milliseconds (Morse, 2019).




Figure 2. PayPal AI Fraud Detection Pipeline: The five-stage process from transaction receipt to final action decision.

At the core of PayPal's detection system are deep learning models, specifically multi-layer neural networks trained on billions of historical transactions. These models learn complex, non-linear relationships between transaction features and fraud outcomes without being programmed with explicit rules. A transaction that looks fine against a simple rule can still be flagged as high-risk when a neural network recognizes its combination of amount, location, device, time, and history matches a known fraud pattern (Morse, 2019). PayPal employs an ensemble approach where multiple model architectures independently produce risk assessments and a composite score is derived from their combined output, reducing the risk of any single model's blind spots causing systematic misclassification. The academic literature confirms that ensemble methods combining random forests, support vector machines, and neural networks produce the strongest fraud detection results across real-world transaction data (Bahoo et al., 2024). A review of 104 ML fraud detection studies backs this up, finding that supervised models consistently perform well on real-world transaction data using metrics like F1-score and AUC-ROC (Rodríguez Barrero, 2024).

PayPal's most technically advanced capability is its use of graph neural networks (GNNs) to identify coordinated fraud rings. While conventional models evaluate accounts and transactions in isolation, GNNs map the relational structure between accounts by treating shared device identifiers, IP addresses, email domains, shipping addresses, and bank account numbers

as edges in a network graph (Alarcon et al., 2021). When the GNN detects an unusually dense cluster of connections, it flags the cluster as a potential fraud ring even if no individual account has yet executed a fraudulent transaction. This network-level perspective is particularly effective against synthetic identity fraud and organized account-takeover campaigns. PayPal’s two-sided platform, holding transaction data on both buyer and seller for every payment, gives its GNN models a richer relational dataset than most single-sided processors can access, amplifying the technique’s effectiveness (Alarcon et al., 2021).

Beyond transaction content, PayPal’s system analyzes behavioral signals that reflect how a user physically interacts with their device. Behavioral biometrics capture typing speed and rhythm, touchscreen pressure and swipe patterns, device orientation, and navigation sequences within the app (DigitalDefynd, 2024). Because these patterns reflect deeply ingrained physical habits, they are extremely difficult for a fraudster to replicate even when armed with a victim’s credentials. Device intelligence complements this by constructing a unique fingerprint for every device used on the platform, incorporating operating system version, screen resolution, installed fonts, and network configuration. A transaction from an unrecognized device automatically elevates the risk score. This mechanism directly reflects a principle from the course: AI systems become more accurate through continued exposure to edge cases, because each anomalous pattern confirmed as fraud becomes training data that sharpens future detection (Houston Community College, 2024).

The detection pipeline operates in real time, computing all 700+ risk features and generating a composite score before the payment processor completes authorization, all within approximately 300 milliseconds (Morse, 2019). Figure 3 illustrates how that score maps to one of three action tiers, covering virtually all transaction outcomes from seamless approval to immediate block.

Risk Tier	Score Range	Approx. Share	System Response
 HIGH	0.70 – 1.00	~1% of transactions	Transaction blocked immediately Fraud investigator alert triggered Account flagged for manual review
 MEDIUM	0.40 – 0.69	~5% of transactions	Step-up authentication required Customer prompted for 2FA / OTP Transaction approved only on confirmation
 LOW	0.00 – 0.39	~94% of transactions	Transaction auto-approved instantly No customer friction introduced Payment completes in <300ms

Figure 3. PayPal Risk Score Decision Tiers: The three-tier scoring framework determining system response for each transaction. Approximate share estimates are conceptual.

Critically, the system is adaptive. Confirmed fraud cases are continuously fed back into the training pipeline, causing models to retrain and update their recognition of emerging fraud patterns without manual rule-writing. This closed-loop architecture allows PayPal to stay ahead of adversarial fraud actors who actively probe and adapt to detection systems (Columbus, 2017).

The shift from rule-based to AI-driven fraud detection produced substantial, quantifiable improvements. Figure 4 translates PayPal’s fraud rate advantage into direct dollar impact at its 2022 transaction scale, Figure 5 compares fraud loss rates across the industry landscape, and Table 1 summarizes broader operational performance gains.




Figure 4. Annual Cost of Fraud Impact: Estimated annual fraud losses for PayPal vs. the industry average, applied to PayPal’s 2022 TPV of \$1.36 trillion. Sources: DigitalDefynd (2024); PayPal Holdings (2023).




Figure 5. Fraud Loss Rate Comparison: PayPal AI-powered vs. industry average vs. pre-AI era estimates (% of TPV). Lower is better.

Metric	Industry Baseline	PayPal with AI
Fraud Loss Rate	~1.00%+ of TPV	~0.32% of TPV

Metric	Industry Baseline	PayPal with AI
Transactions Scored Daily	Rule-based batch	45+ million in real time
Risk Signals per Txn	~10–20 static rules	700+ dynamic features
False Positive Rate	High (manual-heavy)	Substantially reduced
Total Payment Volume	N/A	\$1.36 trillion (2022)
Detection Speed	Hours (batch review)	Milliseconds

Table 1. Operational Performance Metrics: Industry Baseline vs. PayPal with AI.

As Figure 4 demonstrates, the gap between PayPal’s 0.32% fraud rate and the industry average of approximately 1.00% translates to an estimated \$9.25 billion in annually avoided fraud losses at PayPal’s 2022 payment volume (DigitalDefynd, 2024; PayPal Holdings, 2023). To put that in perspective, PayPal’s estimated \$4.35 billion in annual fraud losses is less than one-third of what it would pay at the industry average rate. This performance gap is a direct consequence of the AI system’s ability to detect subtle, multi-signal fraud patterns that rule-based systems cannot identify.

Blocking legitimate transactions is a hidden but significant cost for any payment platform. PayPal’s behavioral and device intelligence models substantially reduced false positives by building individualized behavioral baselines for each user. A large transaction at an unfamiliar merchant might trigger a generic flag, but the model recognizes it as consistent with that customer’s spending history and approves it without any friction for the customer (Columbus, 2017). Research in financial AI confirms that methods achieving this kind of behavioral personalization result in operational cost savings ranging from 6% to 25% of total fraud-related losses (Bahoo et al., 2024).

A rule-based system across 200+ markets would require separate, locally maintained rule sets for each country’s fraud landscape. PayPal’s machine learning models generalize across markets because they learn from the full breadth of global transaction data while remaining sensitive to local behavioral norms (PayPal Holdings, 2023). This allows new markets to be onboarded without rebuilding the detection infrastructure, providing a scalability advantage that compounds over time as the platform expands.

The graph neural network component enables PayPal to act against entire coordinated fraud operations rather than individual accounts. By surfacing the relational structure linking fraudulent accounts through shared infrastructure, investigators can dismantle fraud rings wholesale, preventing the same actors from rotating to new accounts after a single-account closure (Alarcon et al., 2021). This shift from reactive account-level response to proactive network-level interdiction represents a qualitative leap in fraud prevention effectiveness.

Despite its transformative results, PayPal's AI fraud detection implementation has confronted significant challenges across technical, ethical, regulatory, and operational dimensions.

Machine learning models trained on financial transaction data face a fundamental class imbalance problem. Fraudulent transactions represent a tiny fraction of all activity, meaning the training dataset is overwhelmingly composed of legitimate examples. A model trained naively on such data can achieve high overall accuracy simply by predicting that every transaction is legitimate, while still failing to detect the vast majority of actual fraud. Addressing this requires specialized techniques such as oversampling of fraud cases, cost-sensitive loss functions, and precision-recall optimization rather than simple accuracy metrics (Bahoo et al., 2024). Infrastructure latency is an equally pressing constraint: scoring 45 million daily transactions in under 300 milliseconds each requires a distributed computing architecture with near-zero downtime, making system reliability as operationally critical as model accuracy (Morse, 2019).

AI fraud detection systems carry a meaningful risk of encoding and amplifying demographic bias. If historical fraud flagging data reflects patterns in which certain customer groups were disproportionately subjected to scrutiny, a model trained on that history may reproduce those inequities at scale. A customer unfairly declined because their behavioral profile correlates with a high-risk demographic segment faces real harm: a blocked transaction, a frozen account, and the burden of contesting an automated decision whose rationale is invisible to them (DigitalDefynd, 2024). Deep learning models and GNNs are also not easy to interpret. When a transaction is declined, the reasoning may trace back to internal calculations that cannot be explained in plain language. Researchers consistently flag model explainability as one of the biggest open problems in financial AI (Bahoo et al., 2024). A systematic review covering banking, corporate, and capital market contexts finds that interpretability and bias remain persistent challenges with no clear regulatory consensus yet (Yuhertiana & Amin, 2024).

Compliance with data protection and financial regulation across 200+ markets is a continuous engineering and legal challenge. In the European Union, the General Data Protection Regulation (GDPR) restricts how personal behavioral data can be retained, processed, and transferred across borders. In the United States, the Fair Credit Reporting Act and state-level consumer protection statutes impose additional constraints on automated decision-making affecting customers. Designing a fraud detection system that satisfies all of these frameworks simultaneously, without sacrificing the data richness that machine learning requires, demands sustained collaboration between data science, legal, and compliance teams (PayPal Holdings, 2023).

Fraudsters are adaptive adversaries. As PayPal's detection system improves, the actors it targets refine their tactics to mimic legitimate behavioral patterns and probe the system's decision boundaries. This means a model that performs well today cannot be left to run on its own indefinitely, as fraudsters are constantly probing its edges (Columbus, 2017). Internally, building analyst trust in AI-generated risk scores presents its own challenge: investigators must be trained to calibrate their reliance on model outputs appropriately, requiring ongoing

performance reporting and clear escalation procedures for edge cases (Houston Community College, 2024).

PayPal's fraud detection platform is one of the most well-documented examples of machine learning applied in financial services. By replacing static rule-based detection with a layered system of deep learning ensembles, graph neural networks, behavioral biometrics, and adaptive retraining pipelines, PayPal achieved a fraud loss rate of approximately 0.32% of TPV, translating to an estimated \$9.25 billion in avoided losses annually compared to the industry average (DigitalDefynd, 2024; PayPal Holdings, 2023). At \$1.36 trillion in annual payment volume, that gap makes a compelling case for AI investment on its own.

These results line up with what the research says about AI in finance. Studies show that machine learning models beat rule-based systems in fraud classification tasks, and PayPal's numbers validate that at a scale few institutions can match (Bahoo et al., 2024). PayPal's retraining loop also shows the course concept in action: AI systems improve through exposure to edge cases, and every confirmed fraud case feeds back into the model to sharpen future detection (Houston Community College, 2024).

For institutions looking to build similar systems, this case study points to a few key takeaways. First, data quality is the foundation. No model performs above the level of the data it is trained on, so accuracy and volume of labeled fraud cases matter from day one. Second, build explainability into the design early. Global regulations are tightening, and adding interpretability after the fact costs far more than including it from the start (Bahoo et al., 2024). Third, treat fraud detection as an ongoing challenge, not a one-time build. Fraudsters adapt, and the model needs to keep up. Finally, track false positive rates alongside fraud losses. A system that blocks too many legitimate transactions causes its own kind of damage to customer trust and needs to be managed just as carefully.

References

1. Alarcon, C., Chen, Y., Dong, X., Zhao, L., & Luo, C. (2021). Heterogeneous graph neural networks for fraud detection and explanation in supply chain finance. arXiv. <https://arxiv.org/abs/2112.10566>
2. Bahoo, S., Cucculelli, M., Goga, X., & Mondolo, J. (2024). Artificial intelligence in Finance: A comprehensive review through bibliometric and content analysis. SN Business & Economics, 4, Article 23. <https://doi.org/10.1007/s43546-023-00618-x>
3. Columbus, L. (2017, July 10). How PayPal is using AI to master fraud detection. Forbes. <https://www.forbes.com/sites/louiscolumbus/2017/07/10/how-paypal-is-using-ai-to-master-fraud-detection>
4. DigitalDefynd. (2024). PayPal AI case study [2024]: Revolutionizing digital payments. <https://digitaldefynd.com/IQ/paypal-ai-case-study>
5. Morse, J. (2019, March 15). How PayPal uses machine learning to spot fraud before it happens. Mashable. <https://mashable.com/article/paypal-fraud-detection-machine-learning>

6. PayPal Holdings, Inc. (2023). 2022 annual report. <https://investor.pypl.com/annual-reports>
7. Rodríguez Barrero, M. S. (2024). Financial fraud detection through the application of machine learning techniques: A literature review. *Humanities and Social Sciences Communications*, 11(1), Article 1130. <https://doi.org/10.1057/s41599-024-03606-0>
8. Yuhertiana, I., & Amin, A. H. (2024). Artificial intelligence driven approaches for financial fraud detection: A systematic literature review. *KnE Social Sciences*, 2024, 448–468. <https://kneopen.com/KnE-Social/article/view/16551/>